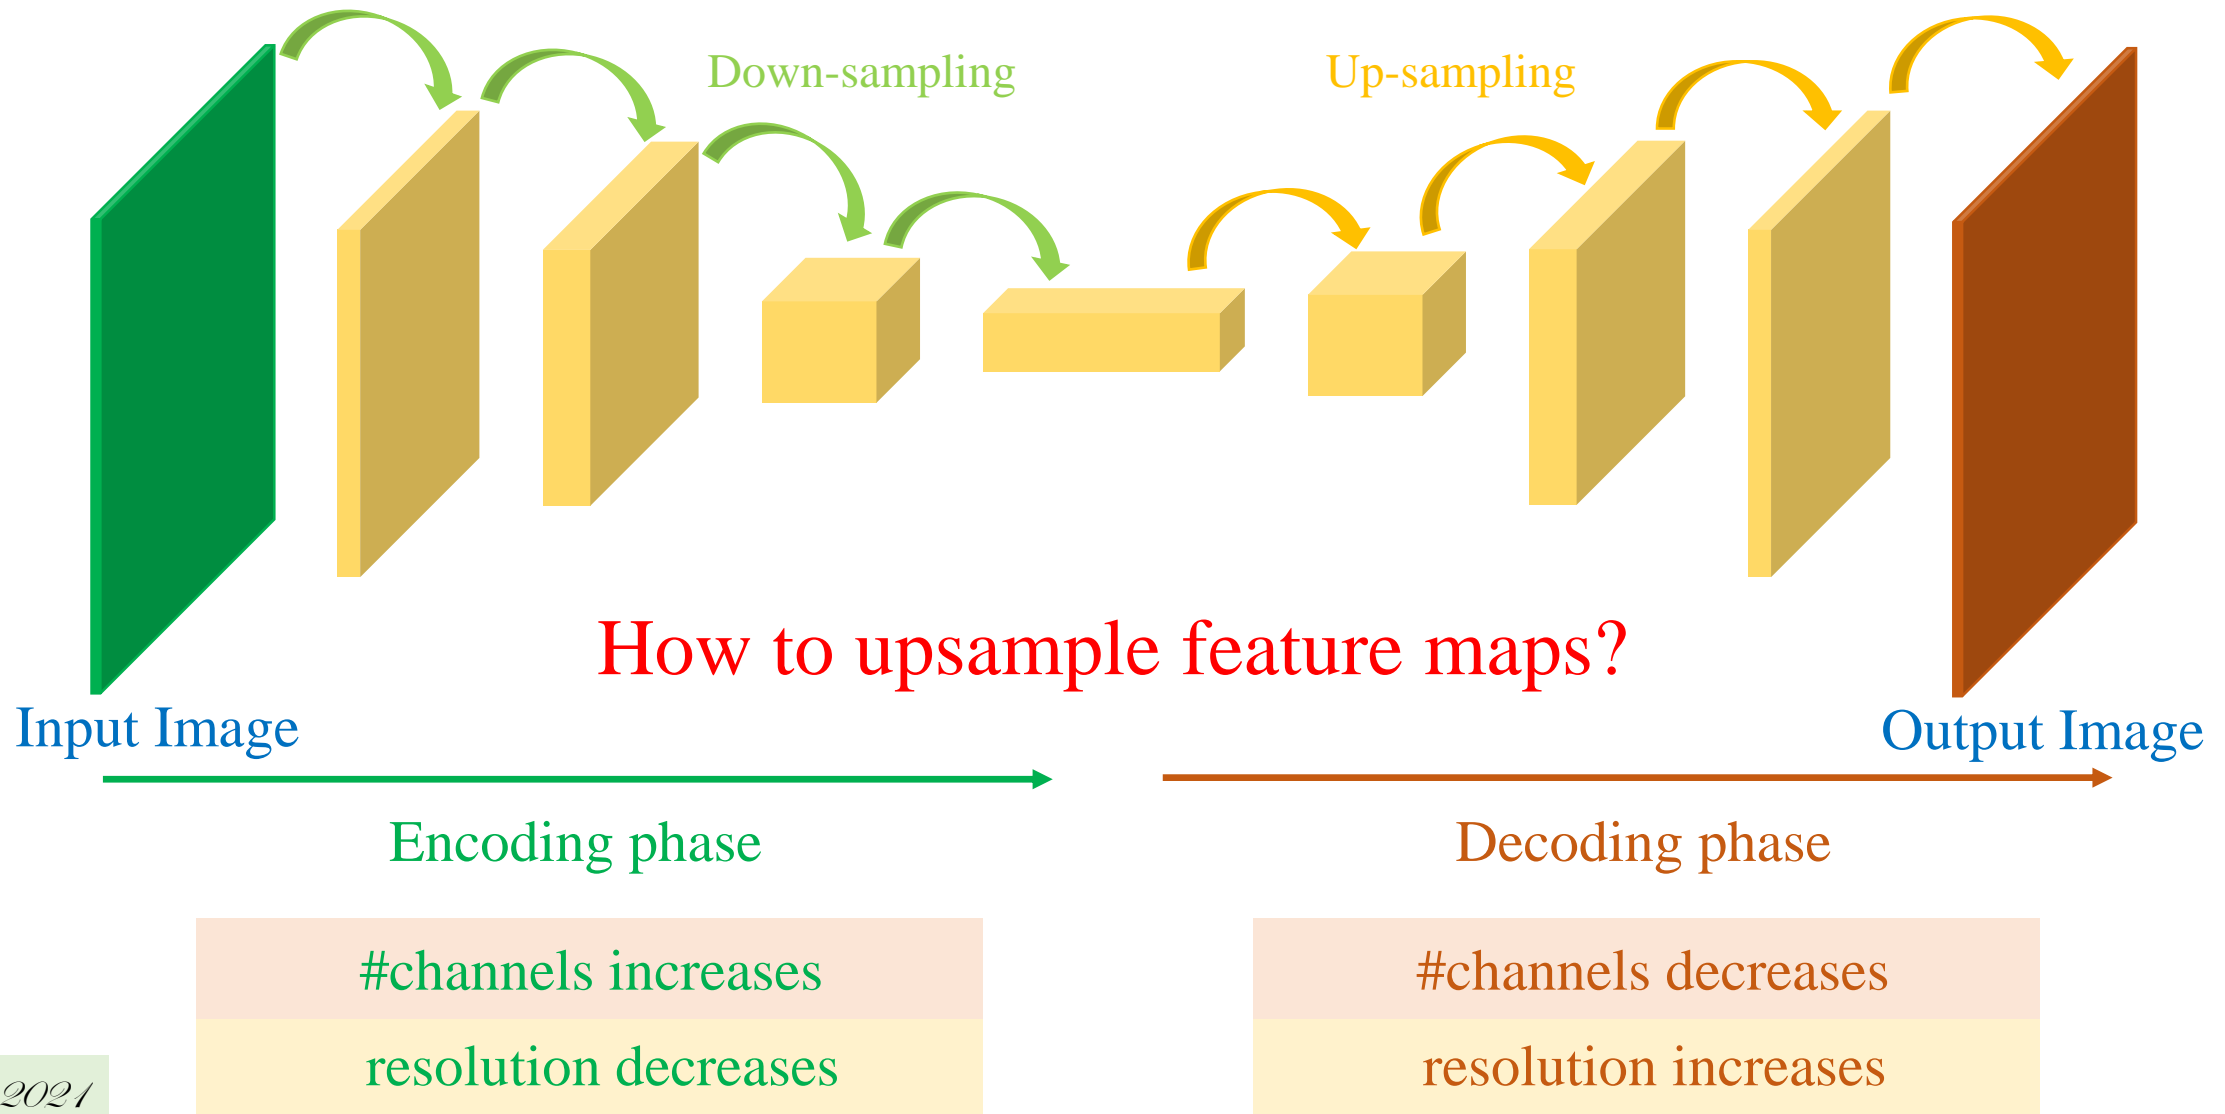


Image Domain Conversion


Quang-Vinh Dinh
Ph.D. in Computer Science

Motivation



How to Upsample Feature Maps


Naïve approach: Only use 'image upsampling'



Output feature maps are lack of details

```
model = tf.keras.Sequential()  
model.add(tf.keras.layers.UpSampling2D(interpolation='bilinear'))  
model.add(tf.keras.layers.Conv2D(num_filters,  
                                kernel_size,  
                                padding='same',  
                                kernel_initializer=initializer))
```

Use 'image upsampling'+Conv



Reduce the weakness
from upsampling

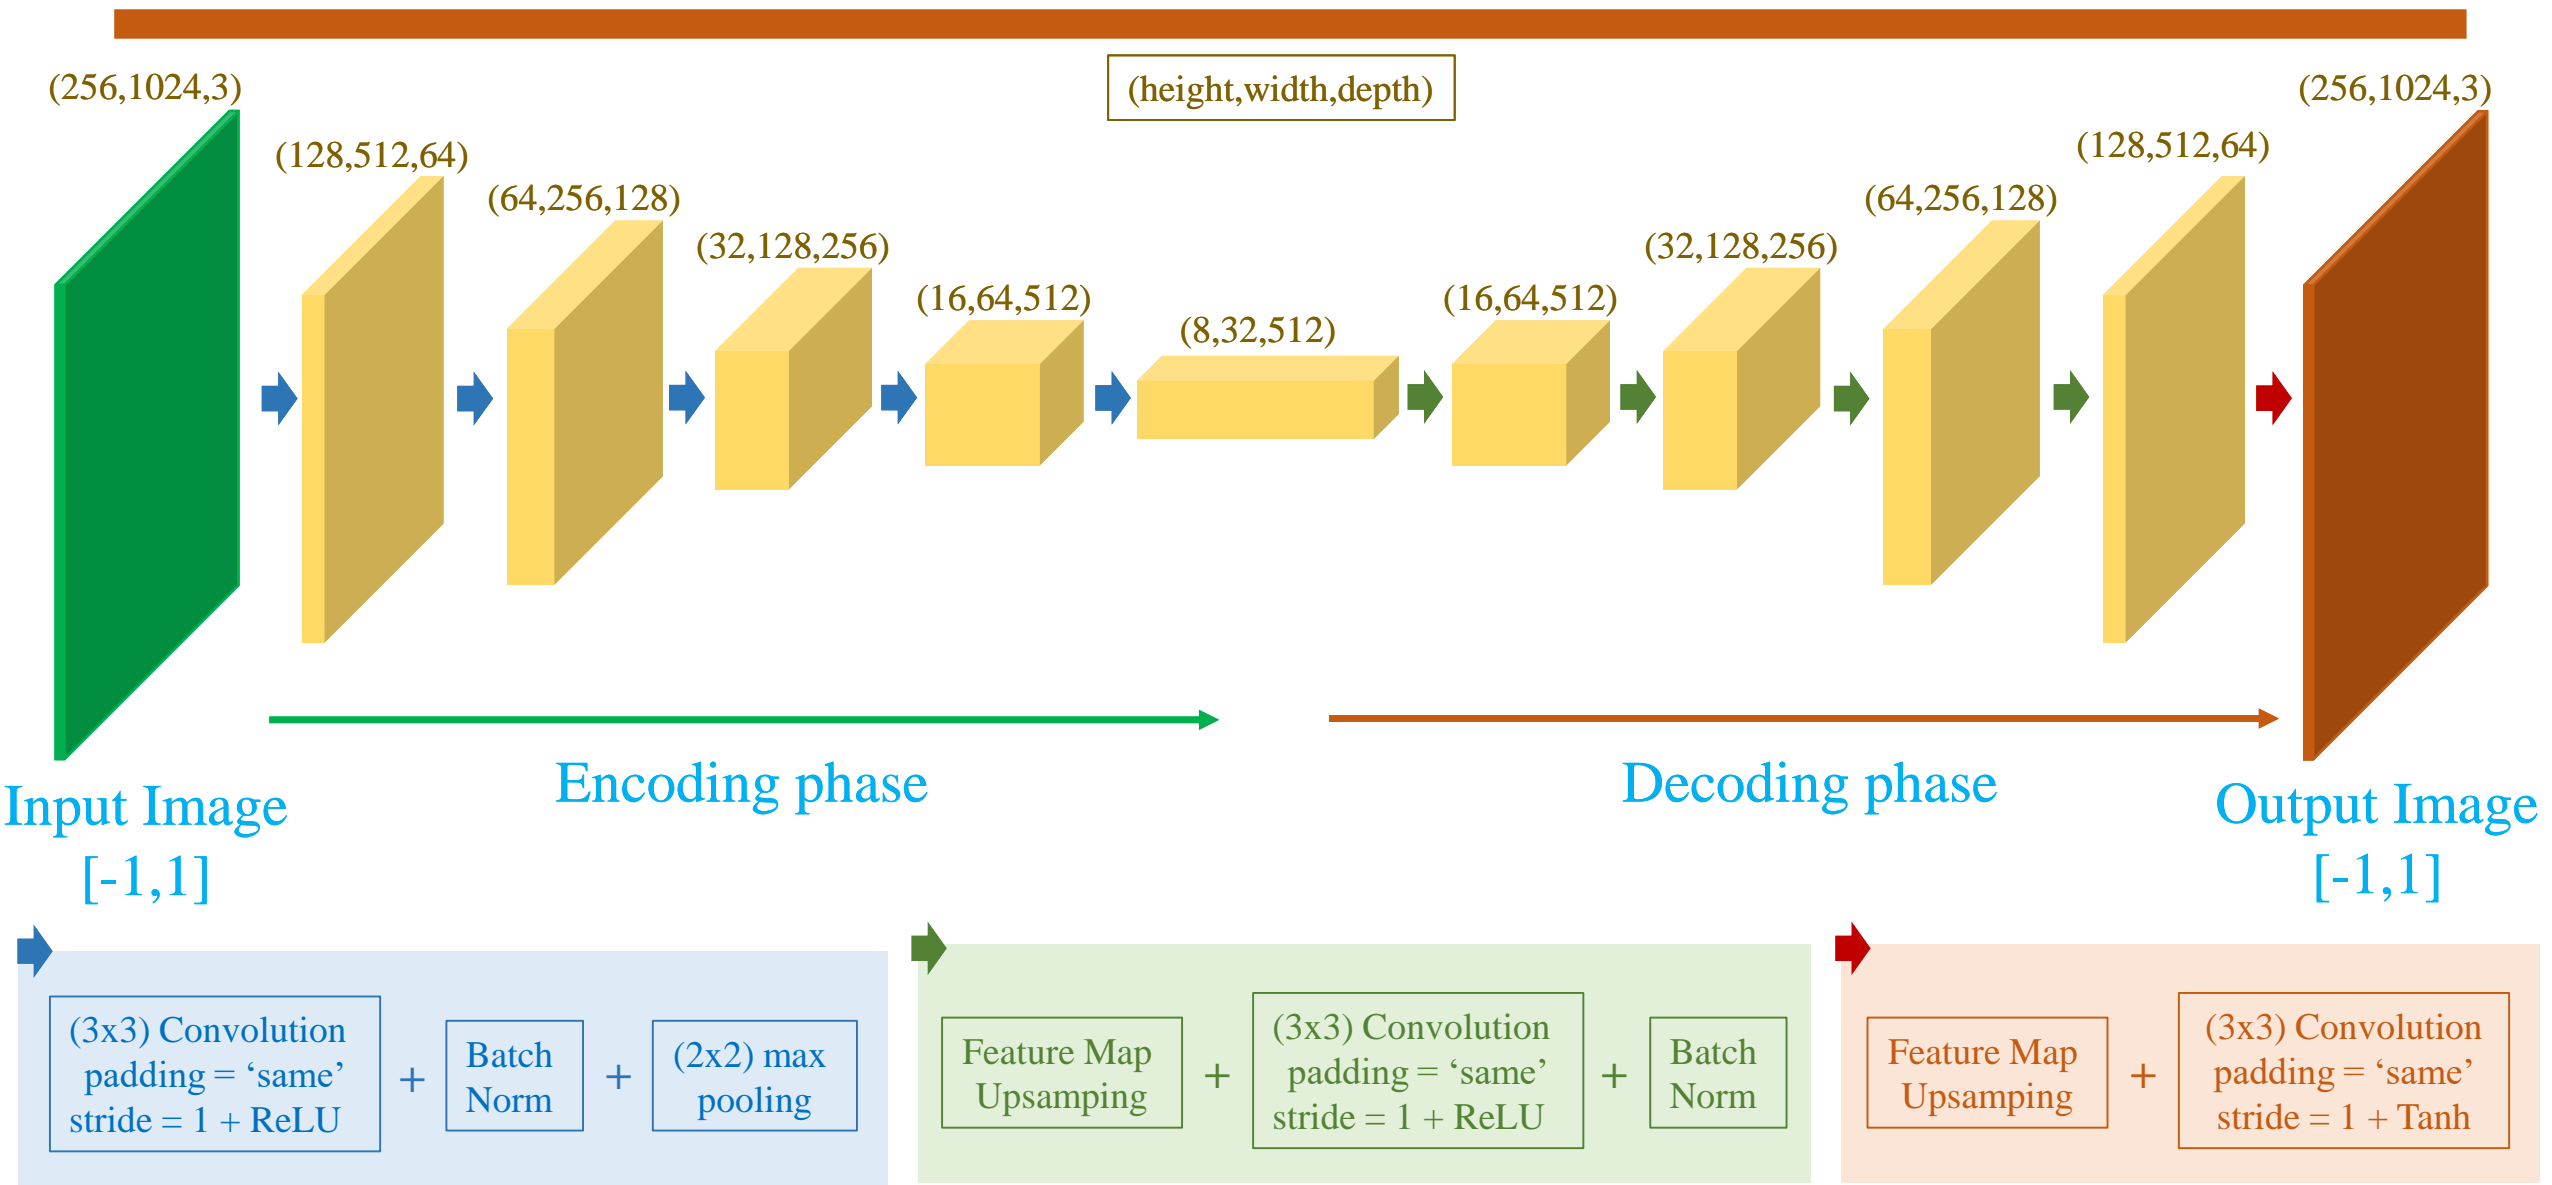


upsampling

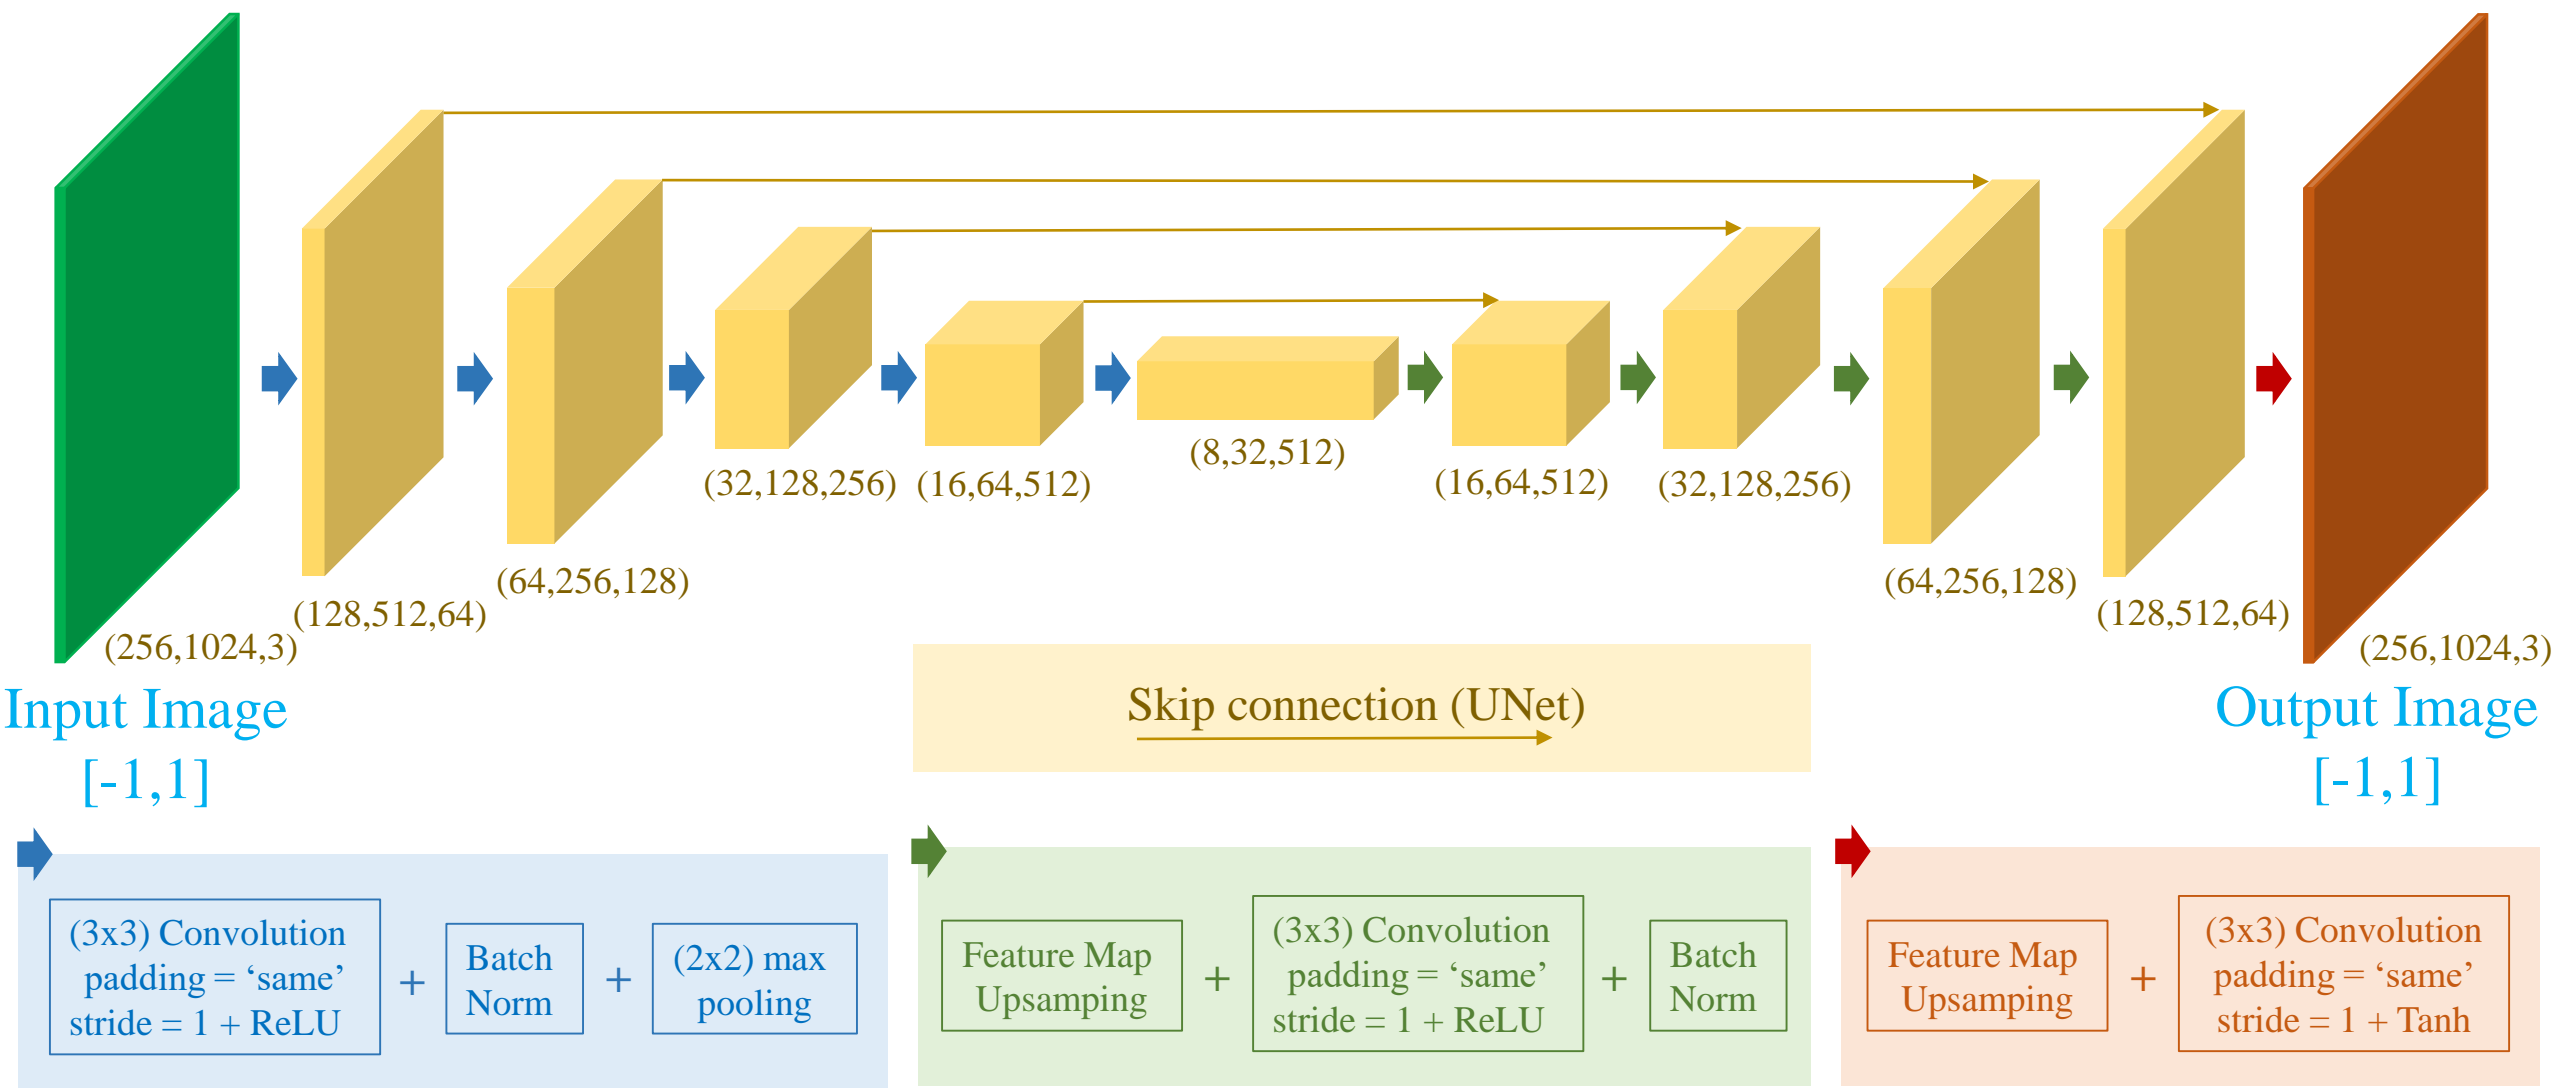


conv
padding='same'
stride=1

How to Upsample Feature Maps




How to Upsample Feature Maps



Applications

❖ Colorization

| | |
|---|----------------------|
| 1 | Data preparation |
| 2 | Network construction |
| 3 | Loss and optimizer |
| 4 | Training |



❖ Colorization

```
PATH = '/content/gdrive/My Drive/data/'
IMG_WIDTH  = 1024
IMG_HEIGHT = 256

def load(image_file):
    image = tf.io.read_file(image_file)
    image = tf.image.decode_jpeg(image)

    # resize
    image = tf.image.resize(image, (IMG_HEIGHT, IMG_WIDTH))

    # copy
    gray = tf.identity(image)
    color = tf.identity(image)

    # convert to gray
    gray = tf.image.rgb_to_grayscale(gray)

    gray = tf.cast(gray, tf.float32)
    color = tf.cast(color, tf.float32)

    return gray, color
```

```
gray, color = load(PATH+'unet/kitti_train/000008_10.png')
```

```
# Show the image
plt.figure()
plt.axis('off')
plt.imshow(gray[:, :, 0]/255.0, cmap='gray')
```

```
plt.figure()
plt.axis('off')
plt.imshow(color/255.0)
```

```
<matplotlib.image.AxesImage at 0x7f7f0324e550>
```



Applications

❖ Colorization

```
BUFFER_SIZE = 50
BATCH_SIZE = 1

# train_dataset
train_dataset = tf.data.Dataset.list_files(PATH+'unet/kitti_train/0000*_10.png')
train_dataset = train_dataset.map(load_image_train,
                                  num_parallel_calls=tf.data.experimental.AUTOTUNE)
train_dataset = train_dataset.shuffle(BUFFER_SIZE)
train_dataset = train_dataset.batch(BATCH_SIZE)

# test_dataset
test_dataset = tf.data.Dataset.list_files(PATH+'unet/kitti_test/0000*_10.png')
test_dataset = test_dataset.map(load_image_test)
test_dataset = test_dataset.batch(1)
```

```
# normalizing the images to [-1, 1]
def normalize(gray, color):
    gray = (gray / 127.5) - 1
    color = (color / 127.5) - 1

    return gray, color

def random_jitter(gray, color):
    if tf.random.uniform(()) > 0.5:
        # random mirroring
        gray = tf.image.flip_left_right(gray)
        color = tf.image.flip_left_right(color)

    return gray, color

def load_image_train(image_file):
    gray, color = load(image_file)
    gray, color = random_jitter(gray, color)
    gray, color = normalize(gray, color)

    return gray, color

def load_image_test(image_file):
    gray, color = load(image_file)
    gray, color = normalize(gray, color)

    return gray, color
```

❖ Network construction

```
def UNet_process(x): # (1, 256, 1024, 3)
    # encoding
    down_stack = [
        downsample(64, 4), # (bs, 128, 512, 64)
        downsample(256, 4), # (bs, 64, 256, 256)
        downsample(512, 4), # (bs, 32, 128, 512)
        downsample(512, 4), # (bs, 16, 64, 512)
        downsample(512, 4), # (bs, 8, 32, 512)
    ]

    skips = []
    for down in down_stack:
        x = down(x)
        skips.append(x)

    skips = reversed(skips[:-1])
```

```
# decoding
up_stack = [
    upsample(512, 4), # (bs, 16, 64, 512)
    upsample(512, 4), # (bs, 32, 128, 512)
    upsample(256, 4), # (bs, 64, 256, 256)
    upsample(64, 4), # (bs, 128, 512, 64)
]

concat = tf.keras.layers.Concatenate()
for up, skip in zip(up_stack, skips):
    x = up(x)
    x = concat([x, skip])

# last layer
OUTPUT_CHANNELS = 3
last = tf.keras.layers.Conv2DTranspose(OUTPUT_CHANNELS,
                                         4, strides=2,
                                         padding='same',
                                         activation='tanh')

x = last(x)

return x
```

Applications

❖ Colorization

3

Loss and optimizer

```
# loss function
def compute_loss(img1, img2):
    return tf.reduce_mean(tf.abs(img1-img2))

#optimizer
optimizer = tf.keras.optimizers.Adam(1e-4, beta_1=0.5)
```

Applications

❖ Colorization

4

Training

```
@tf.function
def train_step(gray, color):
    with tf.GradientTape() as gen_tape:
        # output
        fake_color = generator([gray], training=True)
        loss = compute_loss(fake_color, color)

    gradients = gen_tape.gradient(loss, generator.trainable_variables)
    optimizer.apply_gradients(zip(gradients, generator.trainable_variables))

    return loss

def fit(train_ds, epochs, test_ds):
    for epoch in range(epochs):
        # Train
        for gray, color in train_ds:
            loss = train_step(gray, color)
```

Applications

❖ Colorization

5

For debugging

```
def generate_images(model, gray, real):
    fake = model([gray], training=True)
    plt.figure(figsize=(15,20))

    display_list = [gray[0,:,:,:0], real[0], fake[0]]
    title = ['Input Left', 'Real Left', 'Predicted Left']

    for i in range(3):
        plt.subplot(1, 3, i+1)
        plt.title(title[i])
        plt.imshow(display_list[i] * 0.5 + 0.5)
        plt.axis('off')
    plt.show()
```

```
def evaluate(model, epoch):
    psnr_mean = 0.0
    count = 0
    for gray, real in test_dataset:
        fake = model([gray], training=True)

        psnr = tf.image.psnr(fake*0.5 + 0.5,
                             real*0.5 + 0.5,
                             max_val=1.0)
        psnr = tf.math.reduce_mean(psnr)

        psnr_mean += psnr
        count = count + 1

    psnr_mean = psnr_mean/count
    return psnr_mean
```

Applications

❖ Colorization

4

Training

22.3.UNet-
Colorization_v1_showGrayImage.ipynb

```
@tf.function
def train_step(gray, color):
    with tf.GradientTape() as gen_tape:
        # output
        fake_color = generator([gray], training=True)
        loss = compute_loss(fake_color, color)

    gradients = gen_tape.gradient(loss, generator.trainable_variables)
    optimizer.apply_gradients(zip(gradients, generator.trainable_variables))

    return loss

def fit(train_ds, epochs, test_ds):
    best_pnsr = 0.0
    for epoch in range(epochs):
        # Train
        for gray, color in train_ds:
            loss = train_step(gray, color)

        # for debug
        pnsr = evaluate(generator, epoch)
        if best_pnsr < pnsr:
            best_pnsr = pnsr
            print(best_pnsr)

        for gray, color in test_ds.take(1):
            generate_images(generator, gray, color)
```

Applications

❖ Colorization

5 For debugging

22.3.UNet-Colorization_v3.ipynb

```
def generate_images(model, gray, real):
    fake = model([gray], training=True)
    plt.figure(figsize=(15,20))

    display_list = [gray[0,:,:,:0], real[0], fake[0]]
    title = ['Input Left', 'Real Left', 'Predicted Left']

    i = 0
    plt.subplot(1, 3, i+1)
    plt.title(title[i])
    plt.imshow(display_list[i]*0.5 + 0.5, cmap='gray')
    plt.axis('off')

    for i in range(1,3):
        plt.subplot(1, 3, i+1)
        plt.title(title[i])
        plt.imshow(display_list[i] * 0.5 + 0.5)
        plt.axis('off')
    plt.show()
```


Applications

❖ Denoising

| | |
|---|----------------------|
| 1 | Data preparation |
| 2 | Network construction |
| 3 | Loss and optimizer |
| 4 | Training |

Input Left



Real Left



Predicted Left



Input Left



Real Left



Predicted Left



❖ Denoising

```
def load(image_file):  
    target = tf.io.read_file(image_file)  
    target = tf.image.decode_jpeg(target)  
  
    target = tf.image.resize(target, (IMG_HEIGHT, IMG_WIDTH))  
  
    # make input  
    noise = tf.identity(target)  
  
    noise_level = 30.0  
    n = tf.random.normal((IMG_HEIGHT, IMG_WIDTH, 1))*noise_level  
  
    noise = noise + n  
    noise = tf.clip_by_value(noise, clip_value_min=0, clip_value_max=255)  
  
    return noise, target
```

Applications

❖ Denoising

1

Data preparation

```
noise, target = load(PATH+'unet/kitti_train/000008_10.png')  
  
# casting to int for matplotlib to show the image  
plt.figure()  
plt.axis('off')  
plt.imshow(noise/255.0)  
  
plt.figure()  
plt.axis('off')  
plt.imshow(target/255.0)
```

<matplotlib.image.AxesImage at 0x7f70c0b5d910>



Applications

❖ Denoising

```
BUFFER_SIZE = 50
BATCH_SIZE  = 1

# train_dataset
train_dataset = tf.data.Dataset.list_files(PATH+'unet/kitti_train/0000*_10.png')
train_dataset = train_dataset.map(load_image_train,
                                  num_parallel_calls=tf.data.experimental.AUTOTUNE)
train_dataset = train_dataset.shuffle(BUFFER_SIZE)
train_dataset = train_dataset.batch(BATCH_SIZE)

# test_dataset
test_dataset = tf.data.Dataset.list_files(PATH+'unet/kitti_test/0000*_10.png')
test_dataset = test_dataset.map(load_image_test)
test_dataset = test_dataset.batch(1)
```

Applications

❖ Denoising

```
# normalizing the images to [-1, 1]
def normalize(input_img, target_img):
    input_img = (input_img / 127.5) - 1
    target_img = (target_img / 127.5) - 1

    return input_img, target_img

def random_jitter(input_img, target_img):
    if tf.random.uniform(()) > 0.5:
        # random mirroring
        input_img = tf.image.flip_left_right(input_img)
        target_img = tf.image.flip_left_right(target_img)

    return input_img, target_img

def load_image_train(image_file):
    input_img, target_img = load(image_file)
    input_img, target_img = random_jitter(input_img, target_img)
    input_img, target_img = normalize(input_img, target_img)

    return input_img, target_img

def load_image_test(image_file):
    input_img, target_img = load(image_file)
    input_img, target_img = normalize(input_img, target_img)

    return input_img, target_img
```

❖ Denoising

```
def UNet_process(x): # (1, 256, 1024, 3)
    # encoding
    down_stack = [
        downsample(64, 4), # (bs, 128, 512, 64)
        downsample(256, 4), # (bs, 64, 256, 256)
        downsample(512, 4), # (bs, 32, 128, 512)
        downsample(512, 4), # (bs, 16, 64, 512)
        downsample(512, 4), # (bs, 8, 32, 512)
    ]

    skips = []
    for down in down_stack:
        x = down(x)
        skips.append(x)

    skips = reversed(skips[:-1])
```

```
# decoding
up_stack = [
    upsample(512, 4), # (bs, 16, 64, 512)
    upsample(512, 4), # (bs, 32, 128, 512)
    upsample(256, 4), # (bs, 64, 256, 256)
    upsample(64, 4), # (bs, 128, 512, 64)
]

concat = tf.keras.layers.Concatenate()
for up, skip in zip(up_stack, skips):
    x = up(x)
    x = concat([x, skip])

# last layer
OUTPUT_CHANNELS = 3
last = tf.keras.layers.Conv2DTranspose(OUTPUT_CHANNELS,
                                         4, strides=2,
                                         padding='same',
                                         activation='tanh')

x = last(x)

return x
```


Applications

❖ Denoising

3

Loss and optimizer

```
# loss function
def compute_loss(img1, img2):
    return tf.reduce_mean(tf.abs(img1-img2))

#optimizer
optimizer = tf.keras.optimizers.Adam(1e-4, beta_1=0.5)
```

Applications

❖ Denoising

4

Training

```
@tf.function
def train_step(gray, color):
    with tf.GradientTape() as gen_tape:
        fake_color = generator([gray], training=True)
        loss = compute_loss(fake_color, color)

    gradients = gen_tape.gradient(loss, generator.trainable_variables)
    optimizer.apply_gradients(zip(gradients, generator.trainable_variables))

    return loss

def fit(train_ds, epochs, test_ds):
    for epoch in range(epochs):
        for noise_img, clean_img in train_ds:
            loss = train_step(noise_img, clean_img)
```


Applications

❖ Edge2Scene

1

Data preparation

2

Network construction

3

Loss and optimizer

4

Training

Input Image



Predicted Image



Applications

❖ Edge2Scene

1 Data preparation

```
def load(image_file):  
    image = tf.io.read_file(image_file)  
    image = tf.image.decode_jpeg(image)  
  
    # resize  
    image = tf.image.resize(image, (IMG_HEIGHT, IMG_WIDTH))  
  
    # copy  
    gray = tf.identity(image)  
    color = tf.identity(image)  
  
    # convert to gray  
    gray = tf.image.rgb_to_grayscale(gray)  
    gray = tf.reshape(gray, (1, 256, 1024, 1))  
  
    edges = tf.image.sobel_edges(gray)  
    edges = edges**2  
  
    edges = tf.math.reduce_sum(edges, axis=-1)  
    edges = tf.sqrt(edges)  
  
    edges = tf.cast(edges, tf.float32)  
    color = tf.cast(color, tf.float32)  
  
    return edges[0], color
```

Applications


❖ Edge2Scene

1 Data preparation

```
edges, color = load(PATH+'unet/kitti_train/000008_10.png')
```

```
# Show the image  
plt.figure()  
plt.axis('off')  
plt.imshow(edges[:, :, 0]/255.0, cmap='gray')  
plt.figure()  
plt.axis('off')  
plt.imshow(color/255.0)
```

<matplotlib.image.AxesImage at 0x7f58f019d590>



Applications

❖ Edge2Scene

```
BUFFER_SIZE = 50
BATCH_SIZE  = 1

# train_dataset
train_dataset = tf.data.Dataset.list_files(PATH+'unet/kitti_train/0000*_10.png')
train_dataset = train_dataset.map(load_image_train,
                                  num_parallel_calls=tf.data.experimental.AUTOTUNE)
train_dataset = train_dataset.shuffle(BUFFER_SIZE)
train_dataset = train_dataset.batch(BATCH_SIZE)

# test_dataset
test_dataset = tf.data.Dataset.list_files(PATH+'unet/kitti_test/0000*_10.png')
test_dataset = test_dataset.map(load_image_test)
test_dataset = test_dataset.batch(1)
```

Applications

❖ Edge2Scene

1 Data preparation

```
# normalizing the images to [-1, 1]
def normalize(input_img, target_img):
    input_img = (input_img / 127.5) - 1
    target_img = (target_img / 127.5) - 1

    return input_img, target_img

def random_jitter(input_img, target_img):
    if tf.random.uniform(()) > 0.5:
        # random mirroring
        input_img = tf.image.flip_left_right(input_img)
        target_img = tf.image.flip_left_right(target_img)

    return input_img, target_img

def load_image_train(image_file):
    input_img, target_img = load(image_file)
    input_img, target_img = random_jitter(input_img, target_img)
    input_img, target_img = normalize(input_img, target_img)

    return input_img, target_img

def load_image_test(image_file):
    input_img, target_img = load(image_file)
    input_img, target_img = normalize(input_img, target_img)

    return input_img, target_img
```

❖ Edge2Scene

```
def UNet_process(x): # (1, 256, 1024, 3)
    # encoding
    down_stack = [
        downsample(64, 4), # (bs, 128, 512, 64)
        downsample(256, 4), # (bs, 64, 256, 256)
        downsample(512, 4), # (bs, 32, 128, 512)
        downsample(512, 4), # (bs, 16, 64, 512)
        downsample(512, 4), # (bs, 8, 32, 512)
    ]

    skips = []
    for down in down_stack:
        x = down(x)
        skips.append(x)

    skips = reversed(skips[:-1])
```

```
# decoding
up_stack = [
    upsample(512, 4), # (bs, 16, 64, 512)
    upsample(512, 4), # (bs, 32, 128, 512)
    upsample(256, 4), # (bs, 64, 256, 256)
    upsample(64, 4), # (bs, 128, 512, 64)
]

concat = tf.keras.layers.Concatenate()
for up, skip in zip(up_stack, skips):
    x = up(x)
    x = concat([x, skip])

# last layer
OUTPUT_CHANNELS = 3
last = tf.keras.layers.Conv2DTranspose(OUTPUT_CHANNELS,
                                         4, strides=2,
                                         padding='same',
                                         activation='tanh')

x = last(x)

return x
```


Applications

❖ Edge2Scene

3

Loss and optimizer

```
# loss function
def compute_loss(img1, img2):
    return tf.reduce_mean(tf.abs(img1-img2))

#optimizer
optimizer = tf.keras.optimizers.Adam(1e-4, beta_1=0.5)
```


Applications

❖ Edge2Scene

4

Training

```
@tf.function
def train_step(edge_img, scene_img):
    with tf.GradientTape() as gen_tape:
        fake_scene = generator([edge_img], training=True)
        loss = compute_loss(fake_scene, scene_img)

    gradients = gen_tape.gradient(loss, generator.trainable_variables)
    optimizer.apply_gradients(zip(gradients, generator.trainable_variables))

    return loss

def fit(train_ds, epochs, test_ds):
    best_pnsr = 0.0
    for epoch in range(epochs):
        for edge_img, scene_img in train_ds:
            loss = train_step(edge_img, scene_img)
```


Applications

❖ Super-resolution



1

Data preparation

2

Network construction

3

Loss and optimizer

4

Training

Applications

❖ Super-resolution

1

Data preparation

```
def load(image_file):  
    image = tf.io.read_file(image_file)  
    image = tf.image.decode_jpeg(image)  
  
    image = tf.image.resize(image, (IMG_HEIGHT, IMG_WIDTH))  
  
    # make input  
    lr = tf.identity(image)  
    lr = tf.image.resize(lr, (IMG_HEIGHT//4, IMG_WIDTH//4))  
  
    return lr, image
```

Applications

❖ Super-resolution

1 Data preparation

```
lr, image = load(PATH+'unet/kitti_train/000008_10.png')  
print(lr.shape)  
print(image.shape)
```

```
# casting to int for matplotlib to show the image  
plt.figure()  
plt.axis('off')  
plt.imshow(lr/255.0)  
plt.figure()  
plt.axis('off')  
plt.imshow(image/255.0)
```

(64, 256, 3)

(256, 1024, 3)

<matplotlib.image.AxesImage at 0x7fd2202f77d0>



Applications

❖ Super-resolution

```
BUFFER_SIZE = 50
BATCH_SIZE  = 1

# train_dataset
train_dataset = tf.data.Dataset.list_files(PATH+'unet/kitti_train/0000*_10.png')
train_dataset = train_dataset.map(load_image_train,
                                  num_parallel_calls=tf.data.experimental.AUTOTUNE)
train_dataset = train_dataset.shuffle(BUFFER_SIZE)
train_dataset = train_dataset.batch(BATCH_SIZE)

# test_dataset
test_dataset = tf.data.Dataset.list_files(PATH+'unet/kitti_test/0000*_10.png')
test_dataset = test_dataset.map(load_image_test)
test_dataset = test_dataset.batch(1)
```

Applications

❖ Super-resolution

1 Data preparation

```
# normalizing the images to [-1, 1]
def normalize(input_img, target_img):
    input_img = (input_img / 127.5) - 1
    target_img = (target_img / 127.5) - 1

    return input_img, target_img

def random_jitter(input_img, target_img):
    if tf.random.uniform(()) > 0.5:
        # random mirroring
        input_img = tf.image.flip_left_right(input_img)
        target_img = tf.image.flip_left_right(target_img)

    return input_img, target_img

def load_image_train(image_file):
    input_img, target_img = load(image_file)
    input_img, target_img = random_jitter(input_img, target_img)
    input_img, target_img = normalize(input_img, target_img)

    return input_img, target_img

def load_image_test(image_file):
    input_img, target_img = load(image_file)
    input_img, target_img = normalize(input_img, target_img)

    return input_img, target_img
```

❖ Super-resolution

```
def UNet_process(x): # (1, 256, 1024, 3)
    # encoding
    down_stack = [
        downsample(64, 4), # (bs, 128, 512, 64)
        downsample(256, 4), # (bs, 64, 256, 256)
        downsample(512, 4), # (bs, 32, 128, 512)
        downsample(512, 4), # (bs, 16, 64, 512)
        downsample(512, 4), # (bs, 8, 32, 512)
    ]

    skips = []
    for down in down_stack:
        x = down(x)
        skips.append(x)

    skips = reversed(skips[:-1])
```

```
# decoding
up_stack = [
    upsample(512, 4), # (bs, 16, 64, 512)
    upsample(512, 4), # (bs, 32, 128, 512)
    upsample(256, 4), # (bs, 64, 256, 256)
    upsample(64, 4), # (bs, 128, 512, 64)
]

concat = tf.keras.layers.Concatenate()
for up, skip in zip(up_stack, skips):
    x = up(x)
    x = concat([x, skip])

# last layer
OUTPUT_CHANNELS = 3
last = tf.keras.layers.Conv2DTranspose(OUTPUT_CHANNELS,
                                         4, strides=2,
                                         padding='same',
                                         activation='tanh')

x = last(x)

return x
```


Applications

❖ Super-resolution

3

Loss and optimizer

```
# loss function
def compute_loss(img1, img2):
    return tf.reduce_mean(tf.abs(img1-img2))

#optimizer
optimizer = tf.keras.optimizers.Adam(1e-4, beta_1=0.5)
```

Applications

❖ Super-resolution

4

Training

```
@tf.function
def train_step(lr_img, hr_img):
    with tf.GradientTape() as gen_tape:
        fake_hr = generator([lr_img], training=True)
        loss = compute_loss(fake_hr, hr_img)

        gradients = gen_tape.gradient(loss, generator.trainable_variables)
        optimizer.apply_gradients(zip(gradients, generator.trainable_variables))

    return loss

def fit(train_ds, epochs, test_ds):
    for epoch in range(epochs):
        for lr_img, hr_img in train_ds:
            loss = train_step(lr_img, hr_img)
```


Applications

❖ Deblur

1

Data preparation

2

Network construction

3

Loss and optimizer

4

Training

Input Left



Real Left



Predicted Left



Input Left



Real Left



Predicted Left



Applications

❖ Deblur

1

Data preparation

```
def load(image_file):  
    image = tf.io.read_file(image_file)  
    image = tf.image.decode_jpeg(image)  
  
    image = tf.image.resize(image, (IMG_HEIGHT, IMG_WIDTH))  
  
    # make input  
    blur = tf.identity(image)  
  
    blur = tf.reshape(blur, (1, 256, 1024, 3))  
  
    blur = tfa.image.mean_filter2d(blur)  
    blur = tfa.image.mean_filter2d(blur)  
    blur = tfa.image.mean_filter2d(blur)  
  
    return blur[0], image
```

Applications

❖ Deblur

1 Data preparation

```
blur, target = load(PATH+'unet/kitti_train/000008_10.png')
```

```
# Show the image  
plt.figure()  
plt.axis('off')  
plt.imshow(blur/255.0)  
  
plt.figure()  
plt.axis('off')  
plt.imshow(target/255.0)
```

<matplotlib.image.AxesImage at 0x7f85b437f2d0>



Applications

❖ Deblur

```
BUFFER_SIZE = 50
BATCH_SIZE  = 1

# train_dataset
train_dataset = tf.data.Dataset.list_files(PATH+'unet/kitti_train/0000*_10.png')
train_dataset = train_dataset.map(load_image_train,
                                  num_parallel_calls=tf.data.experimental.AUTOTUNE)
train_dataset = train_dataset.shuffle(BUFFER_SIZE)
train_dataset = train_dataset.batch(BATCH_SIZE)

# test_dataset
test_dataset = tf.data.Dataset.list_files(PATH+'unet/kitti_test/0000*_10.png')
test_dataset = test_dataset.map(load_image_test)
test_dataset = test_dataset.batch(1)
```

Applications

❖ Deblur

1 Data preparation

```
# normalizing the images to [-1, 1]
def normalize(input_img, target_img):
    input_img = (input_img / 127.5) - 1
    target_img = (target_img / 127.5) - 1

    return input_img, target_img

def random_jitter(input_img, target_img):
    if tf.random.uniform(()) > 0.5:
        # random mirroring
        input_img = tf.image.flip_left_right(input_img)
        target_img = tf.image.flip_left_right(target_img)

    return input_img, target_img

def load_image_train(image_file):
    input_img, target_img = load(image_file)
    input_img, target_img = random_jitter(input_img, target_img)
    input_img, target_img = normalize(input_img, target_img)

    return input_img, target_img

def load_image_test(image_file):
    input_img, target_img = load(image_file)
    input_img, target_img = normalize(input_img, target_img)

    return input_img, target_img
```


❖ Deblur

```
def UNet_process(x): # (1, 256, 1024, 3)
    # encoding
    down_stack = [
        downsample(64, 4), # (bs, 128, 512, 64)
        downsample(256, 4), # (bs, 64, 256, 256)
        downsample(512, 4), # (bs, 32, 128, 512)
        downsample(512, 4), # (bs, 16, 64, 512)
        downsample(512, 4), # (bs, 8, 32, 512)
    ]

    skips = []
    for down in down_stack:
        x = down(x)
        skips.append(x)

    skips = reversed(skips[:-1])
```

```
# decoding
up_stack = [
    upsample(512, 4), # (bs, 16, 64, 512)
    upsample(512, 4), # (bs, 32, 128, 512)
    upsample(256, 4), # (bs, 64, 256, 256)
    upsample(64, 4), # (bs, 128, 512, 64)
]

concat = tf.keras.layers.Concatenate()
for up, skip in zip(up_stack, skips):
    x = up(x)
    x = concat([x, skip])

# last layer
OUTPUT_CHANNELS = 3
last = tf.keras.layers.Conv2DTranspose(OUTPUT_CHANNELS,
                                         4, strides=2,
                                         padding='same',
                                         activation='tanh')

x = last(x)

return x
```

Applications

❖ Deblur

3

Loss and optimizer

```
# loss function
def compute_loss(img1, img2):
    return tf.reduce_mean(tf.abs(img1-img2))

#optimizer
optimizer = tf.keras.optimizers.Adam(1e-4, beta_1=0.5)
```

Applications

❖ Deblur

4

Training

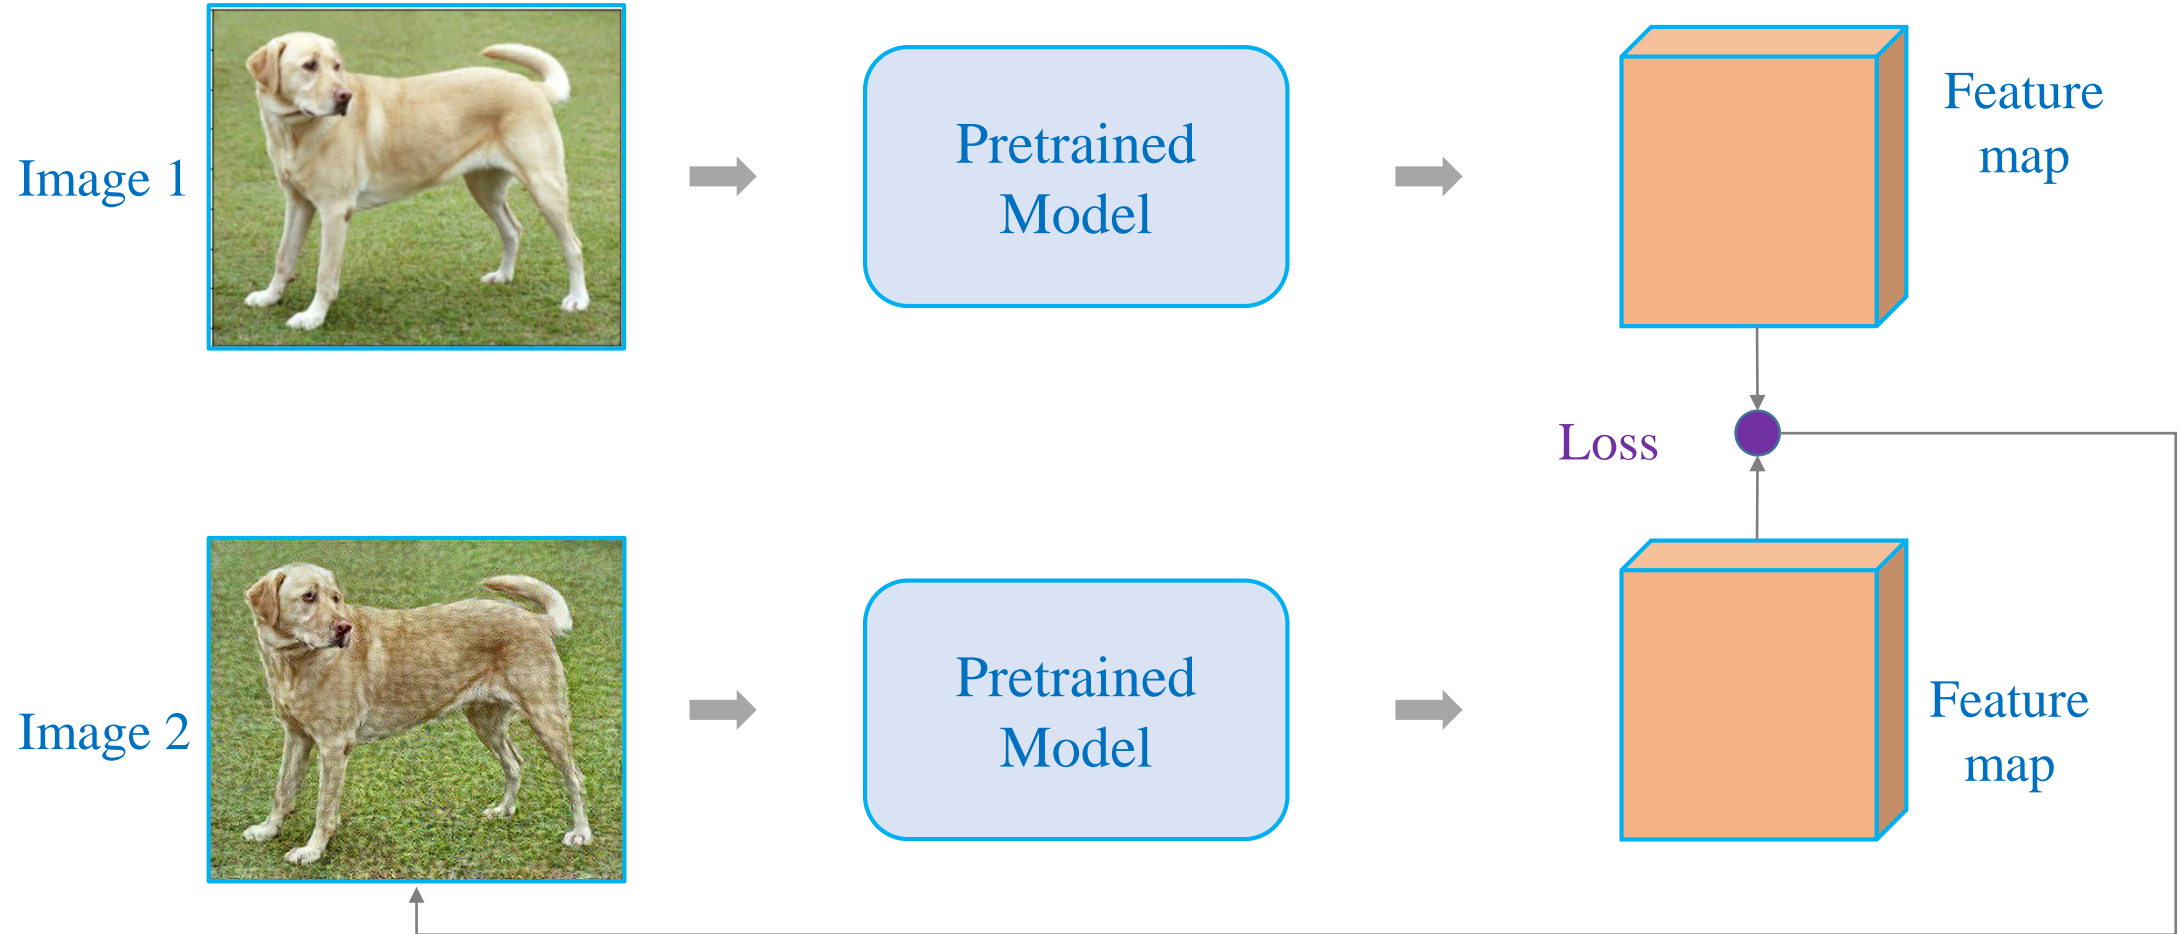
```
@tf.function
def train_step(blur_img, delur_img):
    with tf.GradientTape() as gen_tape:
        fake_delur = generator([blur_img], training=True)
        loss = compute_loss(fake_delur, delur_img)

    gradients = gen_tape.gradient(loss, generator.trainable_variables)
    optimizer.apply_gradients(zip(gradients, generator.trainable_variables))

    return loss

def fit(train_ds, epochs, test_ds):
    best_pnsr = 0.0
    for epoch in range(epochs):
        for blur_img, delur_img in train_ds:
            loss = train_step(blur_img, delur_img)
```


Perceptual Loss



Perceptual Loss

❖ Code

```
from tensorflow.python.keras.applications.vgg19 import VGG19

def get_vgg(output_layer):
    vgg = VGG19(input_shape=(IMG_HEIGHT, IMG_WIDTH, 3),
                include_top=False, weights='imagenet')
    return tf.keras.Model(vgg.input, vgg.layers[20].output)

vgg = get_vgg()
```

23.6.UNet-Colorization-
PerceptualLoss.ipynb

```
img1 = tf.random.normal((1, IMG_HEIGHT, IMG_WIDTH, 3))
img2 = tf.random.normal((1, IMG_HEIGHT, IMG_WIDTH, 3))

img1_fea = vgg(img1)
img2_fea = vgg(img2)

loss = mean_squared_error(img1_fea, img2_fea)
```

Perceptual Loss

❖ Code

```
from tensorflow.python.keras.applications.vgg19 import VGG19

def get_vgg(output_layer):
    vgg = VGG19(input_shape=(IMG_HEIGHT, IMG_WIDTH, 3),
                include_top=False, weights='imagenet')
    return tf.keras.Model(vgg.input, vgg.layers[20].output)
```

```
vgg = get_vgg()
```

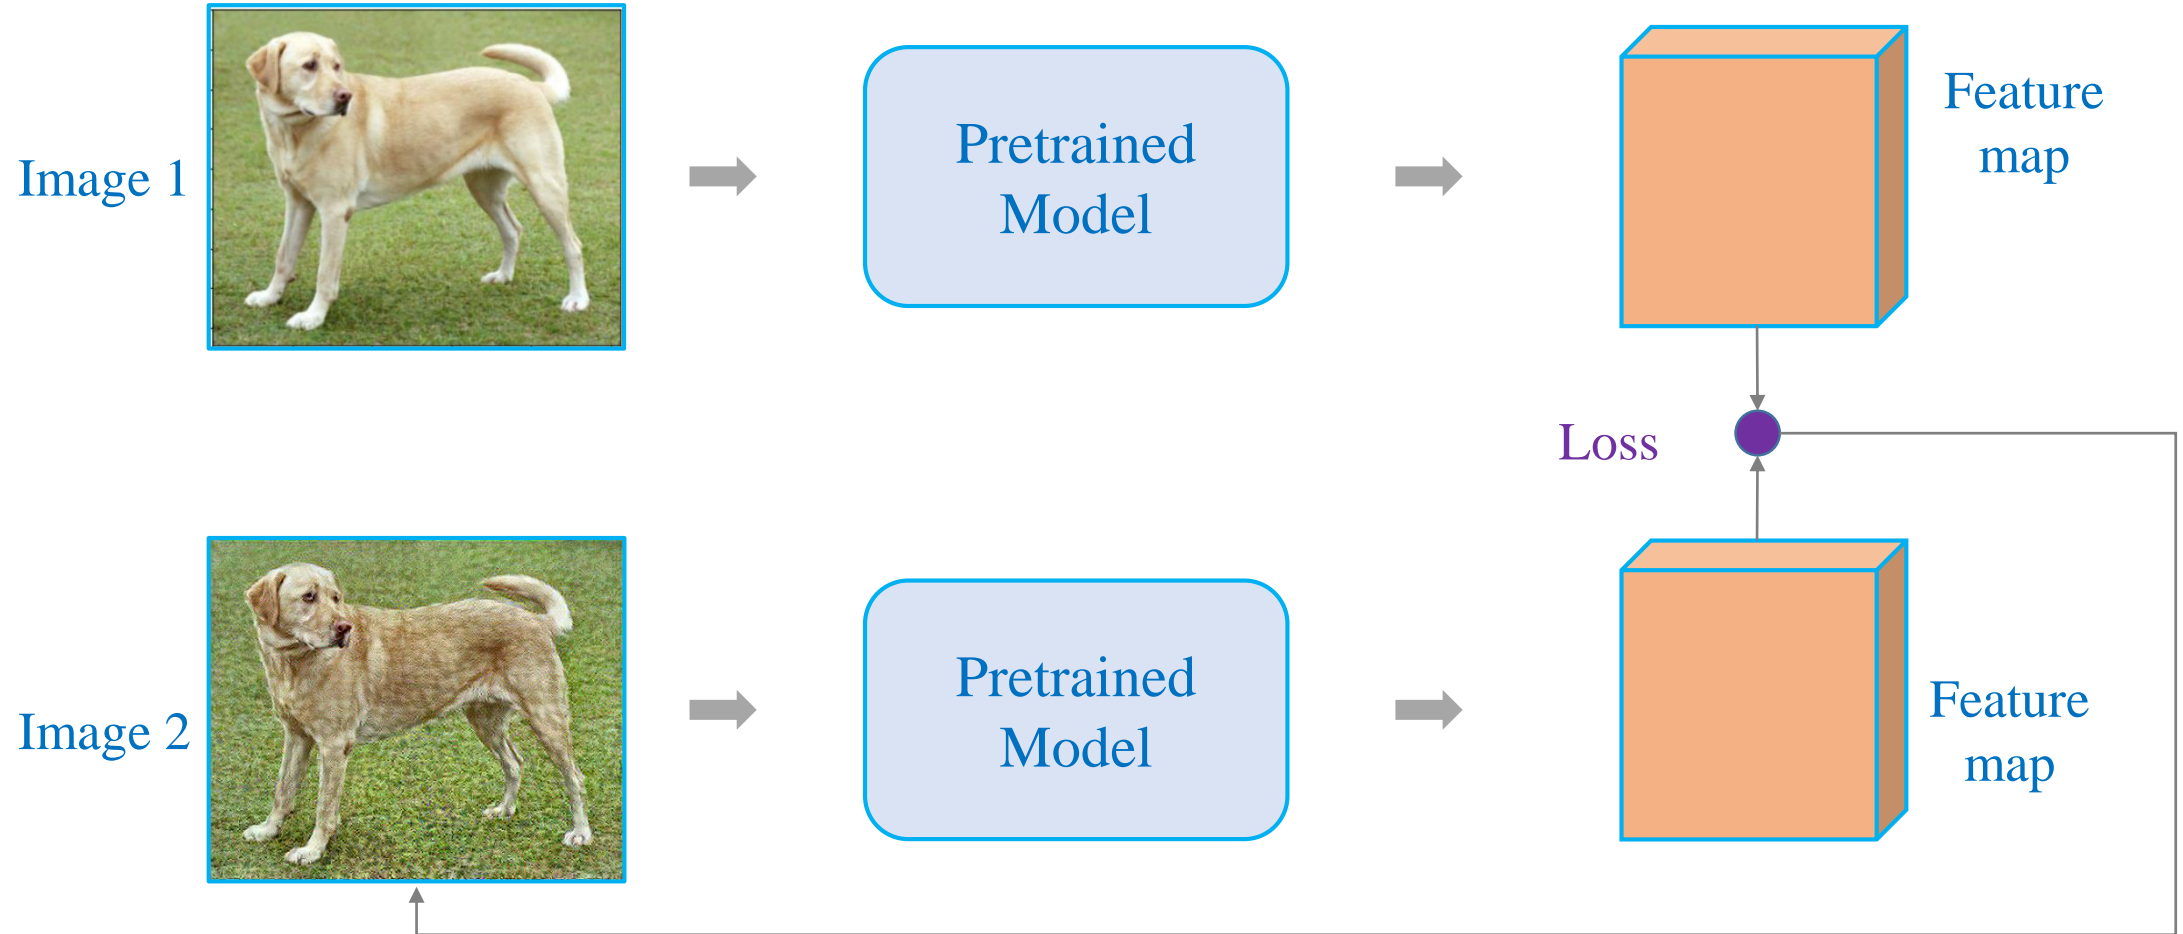
```
mean_squared_error = tf.keras.losses.MeanSquaredError()
```

```
def compute_perceptual_loss(img1, img2):
    if img1.shape[3]==1:
        img1 = tf.concat([img1, img1, img1], 3)
    if img2.shape[3]==1:
        img2 = tf.concat([img2, img2, img2], 3)
    img1_fea = vgg(img1)
    img2_fea = vgg(img2)
```

```
loss = mean_squared_error(img1_fea, img2_fea)
return loss
```

23.6.UNet-Colorization-
PerceptualLoss.ipynb

Perceptual Loss + L1



Perceptual Loss + L1

❖ Code

```
mean_squared_error = tf.keras.losses.MeanSquaredError()
def compute_perceptual_loss(img1, img2):
    if img1.shape[3]==1:
        img1 = tf.concat([img1, img1, img1], 3)
    if img2.shape[3]==1:
        img2 = tf.concat([img2, img2, img2], 3)
    img1_fea = vgg(img1)
    img2_fea = vgg(img2)

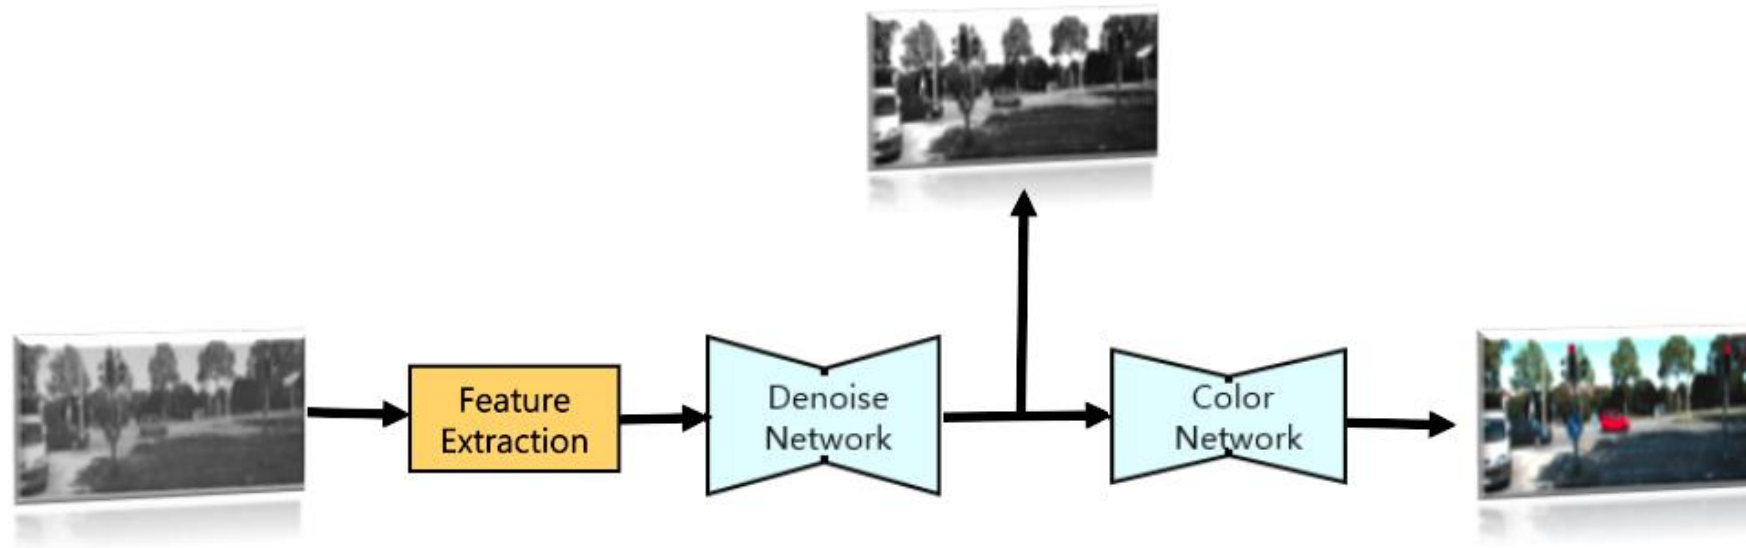
    loss = mean_squared_error(img1_fea, img2_fea)
    return loss
```

```
def compute_loss_l1(img1, img2):
    return tf.reduce_mean(tf.abs(img1-img2))
```

```
def compute_loss(img1, img2):
    return compute_loss_l1(img1, img2) + compute_perceptual_loss(img1, img2)*100.0
```


Denoising+Colorization

❖ Complex domain conversion



Denoising+Colorization

❖ Loss function

Experimental Results

❖ Qualitative result



(a) Input image



(b) Gray result ($PNSR: 31.10$, $SSIM: 0.944$)



(c) Gray target image



(d) Color result ($PNSR: 26.15$, $SSIM: 0.901$)



(e) Color target image

Experimental Results

❖ Compare with a baseline method



Input image



Pix2Pix
(PNSR: 16.49, SSIM: 0.702)



Our
(PNSR: 22.45, SSIM: 0.781)



Input image



Pix2Pix
(PNSR: 18.66, SSIM: 0.805)



Our
(PNSR: 22.72, SSIM: 0.876)



Input image



Pix2Pix
(PNSR: 18.08, SSIM: 0.708)



Our
(PNSR: 24.48, SSIM: 0.831)



Input image



Pix2Pix
(PNSR: 16.28, SSIM: 0.713)



Our
(PNSR: 24.21, SSIM: 0.864)

Ablation Studies

❖ Different losses



(a) Target image



(b) Loss1 ($PNSR: 26.51$, $SSIM: 0.789$)



(c) Loss2 ($PNSR: 25.53$, $SSIM: 0.747$)



(d) Loss3 ($PNSR: 24.91$, $SSIM: 0.706$)



(e) Loss4 ($PNSR: 23.11$, $SSIM: 0.716$)



(f) Loss5 ($PNSR: 22.31$, $SSIM: 0.706$)

